

**Objective:** Make a model to predict the app rating, with other information about the app provided.

**Problem Statement:**

Google Play Store team is about to launch a new feature wherein, certain apps that are promising, are boosted in visibility. The boost will manifest in multiple ways including higher priority in recommendations sections (“Similar apps”, “You might also like”, “New and updated games”). These will also get a boost in search results visibility. This feature will help bring more attention to newer apps that have the potential.

**Domain: General**

Analysis to be done: The problem is to identify the apps that are going to be good for Google to promote. App ratings, which are provided by the customers, is always a great indicator of the goodness of the app. The problem reduces to: predict which apps will have high ratings.

**Content: Dataset:** Google Play Store data (“googleplaystore.csv”)

## **pip install seaborn**

Defaulting to user installation because normal site-packages is not writeable  
Requirement already satisfied: seaborn in c:\programdata\anaconda3\lib\site-packages (0.11.2)

Requirement already satisfied: scipy>=1.0 in

c:\programdata\anaconda3\lib\site-packages (from seaborn) (1.7.3)

Requirement already satisfied: numpy>=1.15 in

c:\programdata\anaconda3\lib\site-packages (from seaborn) (1.21.5)

Requirement already satisfied: matplotlib>=2.2 in

c:\programdata\anaconda3\lib\site-packages (from seaborn) (3.5.1)

Requirement already satisfied: pandas>=0.23 in

c:\programdata\anaconda3\lib\site-packages (from seaborn) (1.4.2)

Requirement already satisfied: cycler>=0.10 in

c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (0.11.0)

Requirement already satisfied: kiwisolver>=1.0.1 in

c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (1.3.2)

Requirement already satisfied: pyparsing>=2.2.1 in

c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (3.0.4)

Requirement already satisfied: fonttools>=4.22.0 in

c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (4.25.0)

Requirement already satisfied: packaging>=20.0 in

c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (21.3)

Requirement already satisfied: pillow>=6.2.0 in

c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (9.0.1)

Requirement already satisfied: python-dateutil>=2.7 in

c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.2->seaborn) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in

c:\programdata\anaconda3\lib\site-packages (from pandas>=0.23->seaborn) (2021.3)

Requirement already satisfied: six>=1.5 in c:\programdata\anaconda3\lib\site-packages (from python-dateutil>=2.7->matplotlib>=2.2->seaborn) (1.16.0)

Note: you may need to restart the kernel to use updated packages.

## **pip install sklearn**

Defaulting to user installation because normal site-packages is not writeable  
Collecting sklearn

Downloading sklearn-0.0.tar.gz (1.1 kB)

Requirement already satisfied: scikit-learn in

c:\programdata\anaconda3\lib\site-packages (from sklearn) (1.0.2)

Requirement already satisfied: scipy>=1.1.0 in

c:\programdata\anaconda3\lib\site-packages (from scikit-learn->sklearn)

```

(1.7.3)
Requirement already satisfied: joblib>=0.11 in
c:\programdata\anaconda3\lib\site-packages (from scikit-learn->sklearn)
(1.1.0)
Requirement already satisfied: numpy>=1.14.6 in
c:\programdata\anaconda3\lib\site-packages (from scikit-learn->sklearn)
(1.21.5)
Requirement already satisfied: threadpoolctl>=2.0.0 in
c:\programdata\anaconda3\lib\site-packages (from scikit-learn->sklearn)
(2.2.0)
Building wheels for collected packages: sklearn
  Building wheel for sklearn (setup.py): started
  Building wheel for sklearn (setup.py): finished with status 'done'
  Created wheel for sklearn: filename=sklearn-0.0-py2.py3-none-any.whl
size=1310
sha256=68ab56c49595bd8a4e5f5dd1e0da972ede244c6884aa46b218264b9f07d5f83b
  Stored in directory:
c:\users\shiwa\appdata\local\pip\cache\wheels\ea\7b\98\b6466d71b8d738a0c54700
8b9eb39bf8676d1ff6ca4b22af1c
Successfully built sklearn
Installing collected packages: sklearn
Successfully installed sklearn-0.0
Note: you may need to restart the kernel to use updated packages.

```

### **pip install matplotlib**

```

Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: matplotlib in
c:\programdata\anaconda3\lib\site-packages (3.5.1)
Requirement already satisfied: packaging>=20.0 in
c:\programdata\anaconda3\lib\site-packages (from matplotlib) (21.3)
Requirement already satisfied: numpy>=1.17 in
c:\programdata\anaconda3\lib\site-packages (from matplotlib) (1.21.5)
Requirement already satisfied: pillow>=6.2.0 in
c:\programdata\anaconda3\lib\site-packages (from matplotlib) (9.0.1)
Requirement already satisfied: python-dateutil>=2.7 in
c:\programdata\anaconda3\lib\site-packages (from matplotlib) (2.8.2)
Requirement already satisfied: cyclor>=0.10 in
c:\programdata\anaconda3\lib\site-packages (from matplotlib) (0.11.0)
Requirement already satisfied: kiwisolver>=1.0.1 in
c:\programdata\anaconda3\lib\site-packages (from matplotlib) (1.3.2)
Requirement already satisfied: fonttools>=4.22.0 in
c:\programdata\anaconda3\lib\site-packages (from matplotlib) (4.25.0)
Requirement already satisfied: pyparsing>=2.2.1 in
c:\programdata\anaconda3\lib\site-packages (from matplotlib) (3.0.4)
Requirement already satisfied: six>=1.5 in c:\programdata\anaconda3\lib\site-
packages (from python-dateutil>=2.7->matplotlib) (1.16.0)
Note: you may need to restart the kernel to use updated packages.

```

- `import numpy as np`  
`import pandas as pd`  
`import matplotlib.pyplot as plt`  
`import seaborn as sns`  
`import statistics as stc`  
`from sklearn.model_selection import train_test_split`  
`from sklearn.linear_model import LinearRegression`  
`from sklearn.metrics import r2_score`  
`from warnings import filterwarnings`  
`filterwarnings('ignore')`  
`plt.rcParams['figure.figsize'] = [15 , 8]`

### Q1. Load the data file using pandas.

- `df = pd.read_csv (r'C:\simplilearn\python project  
simplilearn\googleplaystore.csv')`
- `df.head()`

	App	Category	Rating
0	Photo Editor & Candy Camera & Grid & ScrapBook	ART_AND_DESIGN	4.1
1	Coloring book moana	ART_AND_DESIGN	3.9
2	U Launcher Lite - FREE Live Cool Themes, Hide ...	ART_AND_DESIGN	4.7
3	Sketch - Draw & Paint	ART_AND_DESIGN	4.5
4	Pixel Draw - Number Art Coloring Book	ART_AND_DESIGN	4.3

	Reviews	Size	Installs	Type	Price	Content Rating
0	159	19M	10,000+	Free	0	Everyone
1	967	14M	500,000+	Free	0	Everyone
2	87510	8.7M	5,000,000+	Free	0	Everyone
3	215644	25M	50,000,000+	Free	0	Teen
4	967	2.8M	100,000+	Free	0	Everyone

	Genres	Last Updated	Current Ver
0	Art & Design	January 7, 2018	1.0.0
1	Art & Design;Pretend Play	January 15, 2018	2.0.0
2	Art & Design	August 1, 2018	1.2.4
3	Art & Design	June 8, 2018	Varies with device
4	Art & Design;Creativity	June 20, 2018	1.1

	Android Ver
0	4.0.3 and up
1	4.0.3 and up
2	4.0.3 and up
3	4.2 and up
4	4.4 and up

## 1) Data inspection

- `df.dtypes`

```
App                object
Category           object
Rating             float64
Reviews            object
Size               object
Installs           object
Type               object
Price              object
Content Rating     object
Genres             object
Last Updated       object
Current Ver        object
Android Ver        object
dtype: object
```

- `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10841 entries, 0 to 10840
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   App                   10841 non-null  object
 1   Category              10841 non-null  object
 2   Rating                9367 non-null   float64
 3   Reviews               10841 non-null  object
 4   Size                  10841 non-null  object
 5   Installs              10841 non-null  object
 6   Type                  10840 non-null  object
 7   Price                 10841 non-null  object
 8   Content Rating        10840 non-null  object
 9   Genres                10841 non-null  object
10   Last Updated          10841 non-null  object
11   Current Ver           10833 non-null  object
12   Android Ver           10838 non-null  object
dtypes: float64(1), object(12)
memory usage: 1.1+ MB
```

- `df.shape`

```
(10841, 13)
```

- `df.describe()`

```

                Rating
count  9367.000000
mean    4.193338
std     0.537431
min     1.000000
25%    4.000000
50%    4.300000
75%    4.500000
max    19.000000

```

- `df.describe(include = object)`

```

                App Category Reviews                Size    Installs    Type    Price
\
count    10841    10841    10841                10841    10841    10840    10841
unique    9660         34    6002                462        22         3        93
top      ROBLOX    FAMILY         0  Varies with device  1,000,000+    Free         0
freq         9    1972    596                1695    1579    10039    10040

```

```

                Content Rating Genres    Last Updated                Current Ver    Android Ver
count            10840    10841            10841            10833            10838
unique              6    120            1378            2832             33
top            Everyone    Tools  August 3, 2018  Varies with device  4.1 and up
freq            8714    842            326            1459            2451

```

**Q2. Check for null values in the data. Get the number of null values for each column.**

- `df.isnull().sum()*100/df.shape[0]`

```

App                0.000000
Category           0.000000
Rating            13.596532
Reviews           0.000000
Size              0.000000
Installs          0.000000
Type              0.009224
Price             0.000000
Content Rating    0.009224
Genres            0.000000
Last Updated      0.000000
Current Ver       0.073794
Android Ver       0.027673
dtype: float64

```

**Interpretation: There are null values present in the data. those are Rating = 13.596532, Type = 0.009224, Content Rating = 0.009224, Current Ver = 0.073794, Android Ver = 0.027673,**

### Q3. Drop records with nulls in any of the columns.

- ```
print("Frame Size before : " , df.shape)
df.dropna(subset=['Rating', 'Type', 'Content Rating', 'Current
Ver','Android Ver'],axis=0, inplace=True)
print("Frame Size After : " , df.shape)
df.isnull().sum(axis=0)
```

Frame Size before : (10841, 13)

Frame Size After : (9360, 13)

```
App                0
Category          0
Rating            0
Reviews           0
Size              0
Installs          0
Type              0
Price             0
Content Rating    0
Genres           0
Last Updated     0
Current Ver      0
Android Ver      0
dtype: int64
```

### Q4.1 Variables seem to have incorrect type and inconsistent formatting. You need to fix them:

Q4.1. Size column has sizes in Kb as well as Mb. To analyze, you'll need to convert these to numeric.

Q4.1.1. Extract the numeric value from the column

Q4.1.2. Multiply the value by 1,000, if size is mentioned in Mb

Q4.2. Reviews is a numeric field that is loaded as a string field. Convert it to numeric (int/float).

Q4.3. Installs field is currently stored as string and has values like 1,000,000+. . Q4.3.1. Treat 1,000,000+ as 1,000,000

Q4.3.2. remove '+', ',', from the field, convert it to integer

Q4.4. Price field is a string and has \$ symbol. Remove '\$' sign, and convert it to numeric.

- `df=df[-df['Size'].str.contains('Var')]`  
`df["Size"]`

```
0      19M
1      14M
2      8.7M
3      25M
4      2.8M
...
10833   619k
10834   2.6M
10836    53M
10837   3.6M
10840    19M
```

Name: Size, Length: 7723, dtype: object

#### Q4.1.1. Extract the numeric value from the column

- `df.loc[:, "SizeNum"] = df.Size.str.rstrip("Mk+")`  
`df["SizeNum"]`

```
0      19
1      14
2      8.7
3      25
4      2.8
...
10833   619
10834   2.6
10836    53
10837   3.6
10840    19
```

Name: SizeNum, Length: 7723, dtype: object

- `df.SizeNum = pd.to_numeric(df["SizeNum"])`  
`df.SizeNum.dtype`
- `dtype('float64')`

#### Q4.1.2 Multiply the value by 1,000, if size is mentioned in Mb

- `df['SizeNum']=np.where(df.Size.str.contains('M'),df.SizeNum*1000,`  
`df.SizeNum)`
- `df.Size=df.SizeNum`  
`df.drop('SizeNum',axis=1,inplace=True)`



**Q4.2. Reviews is a numeric field that is loaded as a string field. Convert it to numeric (int/float).**

- `df.Reviews = pd.to_numeric(df.Reviews)`
- `df.Reviews.dtype`
- `dtype('int64')`

**Q4.3. Installs field is currently stored as string and has values like 1,000,000+. Q4.3.2. remove '+', ',' from the field, convert it to integer**

- `df['Installs']=df.Installs.str.replace("1000000+", "1000000")`  
`df['Installs']=df.Installs.str.replace("+", "")`
- `df.Installs=df.Installs.str.replace(",", "")`  
`df.Installs=pd.to_numeric(df.Installs)`  
`df.Installs.dtype`
- `dtype('int64')`
- `df.head()`

|   | App                                               | Category       | Rating |
|---|---------------------------------------------------|----------------|--------|
| 0 | Photo Editor & Candy Camera & Grid & ScrapBook    | ART_AND_DESIGN | 4.1    |
| 1 | Coloring book moana                               | ART_AND_DESIGN | 3.9    |
| 2 | U Launcher Lite – FREE Live Cool Themes, Hide ... | ART_AND_DESIGN | 4.7    |
| 3 | Sketch - Draw & Paint                             | ART_AND_DESIGN | 4.5    |
| 4 | Pixel Draw - Number Art Coloring Book             | ART_AND_DESIGN | 4.3    |

|   | Reviews | Size    | Installs | Type | Price | Content Rating |
|---|---------|---------|----------|------|-------|----------------|
| 0 | 159     | 19000.0 | 10000    | Free | 0     | Everyone       |
| 1 | 967     | 14000.0 | 500000   | Free | 0     | Everyone       |
| 2 | 87510   | 8700.0  | 5000000  | Free | 0     | Everyone       |
| 3 | 215644  | 25000.0 | 50000000 | Free | 0     | Teen           |
| 4 | 967     | 2800.0  | 100000   | Free | 0     | Everyone       |

|   | Genres                    | Last Updated     | Current Ver        |
|---|---------------------------|------------------|--------------------|
| 0 | Art & Design              | January 7, 2018  | 1.0.0              |
| 1 | Art & Design;Pretend Play | January 15, 2018 | 2.0.0              |
| 2 | Art & Design              | August 1, 2018   | 1.2.4              |
| 3 | Art & Design              | June 8, 2018     | Varies with device |
| 4 | Art & Design;Creativity   | June 20, 2018    | 1.1                |

|   | Android Ver  |
|---|--------------|
| 0 | 4.0.3 and up |
| 1 | 4.0.3 and up |
| 2 | 4.0.3 and up |

- 3 4.2 and up
- 4 4.4 and up

**Q4.4. Price field is a string and has \$ symbol. Remove '\$' sign, and convert it to numeric.**

- `df.Price.value_counts()`

```
0          7146
$0.99       105
$2.99       101
$4.99        63
$1.99        53
...
$6.49         1
$1.29         1
$299.99        1
$379.99         1
$1.20         1
```

Name: Price, Length: 68, dtype: int64

- `df['Price'] = df['Price'].str.replace('$', '')`
- `df.Price.value_counts()`

```
0          7146
0.99        105
2.99        101
4.99         63
1.99         53
...
6.49         1
1.29         1
299.99        1
379.99         1
1.20         1
```

Name: Price, Length: 68, dtype: int64

- `df['Price'] = df['Price'].astype(float)`

## 5. Sanity checks:

5.1. Average rating should be between 1 and 5 as only these values are allowed on the play store.

Drop the rows that have a value outside this range.

5.2. Reviews should not be more than installs as only those who installed can review the app.

If there are any such records, drop them.

5.3. For free apps (type = "Free"), the price should not be >0. Drop any such rows.

### 5.1. Average rating should be between 1 and 5 as only these values are allowed on the play store. Drop the rows that have a value outside this range.

- ```
df=df[(df.Rating>=1) & (df.Rating<=5) ]  
df.head()  
df.tail()
```

	App	Category \
10833	Chemin (fr)	BOOKS_AND_REFERENCE
10834	FR Calculator	FAMILY
10836	Sya9a Maroc - FR	FAMILY
10837	Fr. Mike Schmitz Audio Teachings	FAMILY
10840	iHoroscope - 2018 Daily Horoscope & Astrology	LIFESTYLE

	Rating	Reviews	Size	Installs	Type	Price	Content Rating \
10833	4.8	44	619.0	1000	Free	0.0	Everyone
10834	4.0	7	2600.0	500	Free	0.0	Everyone
10836	4.5	38	53000.0	5000	Free	0.0	Everyone
10837	5.0	4	3600.0	100	Free	0.0	Everyone
10840	4.5	398307	19000.0	10000000	Free	0.0	Everyone

	Genres	Last Updated	Current Ver \
10833	Books & Reference	March 23, 2014	0.8
10834	Education	June 18, 2017	1.0.0
10836	Education	July 25, 2017	1.48
10837	Education	July 6, 2018	1.0
10840	Lifestyle	July 25, 2018	Varies with device

	Android Ver
10833	2.2 and up
10834	4.1 and up
10836	4.1 and up
10837	4.1 and up
10840	Varies with device

**5.2. Reviews should not be more than installs as only those who installed can review the app. If there are any such records, drop them.**

- `len(df.index)`

**7723**

- `df.drop(df.index[df.Reviews>df.Installs],axis=0,inplace=True)`  
`len(df.index)`

**7717**

- **For free apps (type = "Free"), the price should not be >0. Drop any such rows.**
- ```
index_free_and_price_gt_0=df.index[((df.Type=='Free')&(df.Price>0))]  
if len(index_free_and_price_gt_0)>0:  
    print("Dropping following indices:",index_free_and_price_gt_0)  
    df.drop(index_free_and_price_gt_0,axis=0,inplace=True)  
else:  
    print("There is no Free Apps with price >0")
```
- There is no Free Apps with price >0

**Interpretation: There is no Free Apps with price >0**

## 5. Performing univariate analysis:

### 5.1. Boxplot for Price

5.1.1 Are there any outliers? Think about the price of usual apps on Play Store.

### 5.2. Boxplot for Reviews

5.2.1 Are there any apps with very high number of reviews? Do the values seem right?

### 5.3. Histogram for Rating


5.3.1 How are the ratings distributed? Is it more toward higher ratings?

### 5.4. Histogram for Size

5.4.1 Note down your observations for the plots made above. Which of these seem to have outliers?

#### 5.1. Boxplot for Price

- `bprice = sns.boxplot(x='Price', data=df)`



**Interpretation: Most of Price values are less than 50 while there is some near concentration around 80. greater than 100 may be considered outliers**

- `price_standard_deviation=stc.stdev(df.Price)`  
`price_standard_deviation`

17.414783874309933

- `price_mean=stc.mean(df.Price)`  
`price_mean`

1.128724893093171

- `price_outlier_uplimit=price_mean+3*price_standard_deviation`  
`price_outlier_uplimit`

53.37307651602297

- `print("Number of upper outliers is  
",len(df[(df.Price>price_outlier_uplimit) ]))`


Number of upper outliers is 17

**Interpretation: Number of upper outliers is 17**

## 5.2. Boxplot for Reviews

```
sns.boxplot(x='Reviews',data=df)
```

```
<AxesSubplot:xlabel='Reviews'>
```



**Interpretation: Most Apps get about less than 2M review. Roughly, greater than 2M can be considered outliers**

- `review_standard_deviation=stc.stdev(df.Reviews)`  
`review_standard_deviation`

1864639.6094670836

- `review_mean=stc.mean(df.Reviews)`  
`review_mean`

295127.5482700531

- `review_outlier_uplimit=review_mean+3*review_standard_deviation`  
`rev_outlier_uplimit`

5889046.376671304

- `review_outlier_downlimit=review_mean-3*review_standard_deviation`  
`review_outlier_downlimit`

-5298791.280131198

- `print("number of upper outliers is  
",len(df[(df.Reviews>rev_outlier_uplimit) ]))`


number of upper outliers is 89

**Interpretation: number of upper outliers is 89**

### 5.3. Histogram for Rating

```
sns.histplot(x='Rating',data=df)
```

```
<AxesSubplot:xlabel='Rating', ylabel='Count'>
```




### 5.3.1 How are the ratings distributed? Is it more toward higher ratings?

**Interpretation: ratings distributed towards higher rating**


### 5.4. Histogram for Size

```
sns.histplot(x='Size', data=df, log_scale=True)
```

```
<AxesSubplot:xlabel='Size', ylabel='Count'>
```



```
bsize = sns.boxplot(x='Size', data=df)
```





#### 5.4.1 Note down your observations for the plots made above.

**interpretation: most of the app size lies under 100000**

- `Size_standard_deviation=stc.stdev(df.Size)`  
`review_standard_deviation`

1864639.6094670836

- `Size_mean=stc.mean(df.Size)`  
`Size_mean`

22976.614293119088

- `Size_outlier_uplimit=Size_mean+3*Size_standard_deviation`  
`Size_outlier_uplimit`

93346.9260933562

- `Size_outlier_downlimit=Size_mean-3*Size_standard_deviation`  
`Size_outlier_downlimit`

-47393.697507118035

- `print("number of upper outliers is`  
`",len(df[(df.Size>Size_outlier_uplimit) ]))`

**number of upper outliers is 148**

**5.4.1 Note down your observations for the plots made above. Which of these seem to have outliers?**

**interpretation: App which are more than size of 93346.92 consider as outlier.  
number of outlier are 148**

## 6. Outlier treatment:

6.1. Price: From the box plot, it seems like there are some apps with very high price. A price of \$200 for an application on the Play Store is very high and suspicious!

6.1.1. Check out the records with very high price

6.1.1.1. Is 200 indeed a high price?

6.1.2 Drop these as most seem to be junk apps

6.2. Reviews: Very few apps have very high number of reviews. These are all star apps that don't help with the analysis and, in fact, will skew it. Drop records having more than 2 million reviews.

6.3 Installs: There seems to be some outliers in this field too. Apps having very high number of installs should be dropped from the analysis.

6.3.1 Find out the different percentiles - 10, 25, 50, 70, 90, 95, 99

6.3.2 Decide a threshold as cutoff for outlier and drop records having values more than that

6.1 Price: From the box plot, it seems like there are some apps with very high price. A price of \$200 for an application on the Play Store is very high and suspicious!

### 6.1.1 Check out the records with very high price

- `df[df.Price>=200]`

|      | App                            | Category  | Rating | Reviews | Size    | \ |
|------|--------------------------------|-----------|--------|---------|---------|---|
| 4197 | most expensive app (H)         | FAMILY    | 4.3    | 6       | 1500.0  |   |
| 4362 | I'm rich                       | LIFESTYLE | 3.8    | 718     | 26000.0 |   |
| 4367 | I'm Rich - Trump Edition       | LIFESTYLE | 3.6    | 275     | 7300.0  |   |
| 5351 | I am rich                      | LIFESTYLE | 3.8    | 3547    | 1800.0  |   |
| 5354 | I am Rich Plus                 | FAMILY    | 4.0    | 856     | 8700.0  |   |
| 5355 | I am rich VIP                  | LIFESTYLE | 3.8    | 411     | 2600.0  |   |
| 5356 | I Am Rich Premium              | FINANCE   | 4.1    | 1867    | 4700.0  |   |
| 5357 | I am extremely Rich            | LIFESTYLE | 2.9    | 41      | 2900.0  |   |
| 5358 | I am Rich!                     | FINANCE   | 3.8    | 93      | 22000.0 |   |
| 5359 | I am rich(premium)             | FINANCE   | 3.5    | 472     | 965.0   |   |
| 5362 | I Am Rich Pro                  | FAMILY    | 4.4    | 201     | 2700.0  |   |
| 5364 | I am rich (Most expensive app) | FINANCE   | 4.1    | 129     | 2700.0  |   |
| 5366 | I Am Rich                      | FAMILY    | 3.6    | 217     | 4900.0  |   |
| 5369 | I am Rich                      | FINANCE   | 4.3    | 180     | 3800.0  |   |
| 5373 | I AM RICH PRO PLUS             | FINANCE   | 4.0    | 36      | 41000.0 |   |

|      | Installs | Type | Price  | Content | Rating   | Genres        | Last Updated      |
|------|----------|------|--------|---------|----------|---------------|-------------------|
| \    |          |      |        |         |          |               |                   |
| 4197 | 100      | Paid | 399.99 |         | Everyone | Entertainment | July 16, 2018     |
| 4362 | 10000    | Paid | 399.99 |         | Everyone | Lifestyle     | March 11, 2018    |
| 4367 | 10000    | Paid | 400.00 |         | Everyone | Lifestyle     | May 3, 2018       |
| 5351 | 100000   | Paid | 399.99 |         | Everyone | Lifestyle     | January 12, 2018  |
| 5354 | 10000    | Paid | 399.99 |         | Everyone | Entertainment | May 19, 2018      |
| 5355 | 10000    | Paid | 299.99 |         | Everyone | Lifestyle     | July 21, 2018     |
| 5356 | 50000    | Paid | 399.99 |         | Everyone | Finance       | November 12, 2017 |
| 5357 | 1000     | Paid | 379.99 |         | Everyone | Lifestyle     | July 1, 2018      |
| 5358 | 1000     | Paid | 399.99 |         | Everyone | Finance       | December 11, 2017 |
| 5359 | 5000     | Paid | 399.99 |         | Everyone | Finance       | May 1, 2017       |
| 5362 | 5000     | Paid | 399.99 |         | Everyone | Entertainment | May 30, 2017      |
| 5364 | 1000     | Paid | 399.99 |         | Teen     | Finance       | December 6, 2017  |
| 5366 | 10000    | Paid | 389.99 |         | Everyone | Entertainment | June 22, 2018     |
| 5369 | 5000     | Paid | 399.99 |         | Everyone | Finance       | March 22, 2018    |
| 5373 | 1000     | Paid | 399.99 |         | Everyone | Finance       | June 25, 2018     |

|      | Current Ver | Android Ver  |
|------|-------------|--------------|
| 4197 | 1.0         | 7.0 and up   |
| 4362 | 1.0.0       | 4.4 and up   |
| 4367 | 1.0.1       | 4.1 and up   |
| 5351 | 2.0         | 4.0.3 and up |
| 5354 | 3.0         | 4.4 and up   |
| 5355 | 1.1.1       | 4.3 and up   |
| 5356 | 1.6         | 4.0 and up   |
| 5357 | 1.0         | 4.0 and up   |
| 5358 | 1.0         | 4.1 and up   |
| 5359 | 3.4         | 4.4 and up   |

|      |       |              |
|------|-------|--------------|
| 5362 | 1.54  | 1.6 and up   |
| 5364 | 2     | 4.0.3 and up |
| 5366 | 1.5   | 4.2 and up   |
| 5369 | 1.0   | 4.2 and up   |
| 5373 | 1.0.2 | 4.1 and up   |

- `print("Number of Apps with price >= 200 are ",len(df[(df.Price>=200)]))`

**Number of Apps with price >= 200 are 15**

- Is 200 indeed a high price?
- 6.1.2 Drop these as most seem to be junk apps
- `df.drop(df.index[(df.Price>=200)], inplace=True)`  
`len(df.index)`

**7702**

## 6.2. Reviews: Very few apps have very high number of reviews

These are all star apps that don't help with the analysis and, in fact, will skew it. Drop records having more than 2 million reviews.

- `df.drop(df.index[(df.Reviews>=2000000)], inplace=True)`  
`len(df.index)`

**7483**

- 6.3 Installs: There seems to be some outliers in this field too. Apps having very high number of installs should be dropped from the analysis.

### 6.3.1 Find out the different percentiles – 10, 25, 50, 70, 90, 95, 99

```
install_10_perc=np.percentile(df.Installs, 10)
install_10_perc
```

**1000.0**

```
install_25_perc=np.percentile(df.Installs, 25)
install_25_perc
```

**10000.0**

```
install_50_perc=np.percentile(df.Installs, 50)
install_50_perc
```

**100000.0**

```
install_70_perc=np.percentile(df.Installs, 70)
install_70_perc
```

**1000000.0**

```
install_90_perc=np.percentile(df.Installs, 90)  
install_90_perc
```

**10000000.0**

```
install_95_perc=np.percentile(df.Installs, 95)  
install_95_perc
```


**10000000.0**

```
install_99_perc=np.percentile(df.Installs, 99)  
install_99_perc
```

**50000000.0**

```
sns.histplot(data=df,x='Installs',log_scale=True)
```

```
<AxesSubplot:xlabel='Installs', ylabel='Count'>
```



- `Installs_standard_deviation=stc.stdev(df.Installs)`  
`Installs_standard_deviation`

**27818305.317482274**

- `Installs_mean=stc.mean(df.Installs)`  
`Installs_mean`

**3947464.5449685953**

- `Installs_outlier_uplimit=Installs_mean+3*Installs_standard_deviation`  
`Installs_outlier_uplimit`

**87402380.49741541**

- `print("number of upper outliers is`  
`",len(df[(df.Installs>Installs_outlier_uplimit) ]))`

**number of upper outliers is 60**

- `df.drop(df.index[df.Installs >= 87402380.49741541],inplace=True)`  
`len(df.index)`

**7423**

**I Decide more than that of 99% as an threshold cutoff for outlier and drop records having values more than that.**

- `df.drop(df.index[df.Installs >= install_99_perc],inplace=True)`  
`len(df.index)`

**7307**

**7. Bivariate analysis: Let's look at how the available predictors relate to the variable of interest, i.e., our target variable rating. Make scatter plots (for numeric features) and box plots (for character features) to assess the relations between rating and the other features.**

7.1 Make scatter plot/joinplot for Rating vs. Price

7.1.1 What pattern do you observe? Does rating increase with price?

7.2 Make scatter plot/joinplot for Rating vs. Size

7.2.1 Are heavier apps rated better?

7.3 Make scatter plot/joinplot for Rating vs. Reviews

7.3.1 Does more review mean a better rating always?

7.4 Make boxplot for Rating vs. Content Rating

7.4.1 Is there any difference in the ratings? Are some types liked better?

7.5 Make boxplot for Ratings vs. Category


7.5.1 Which genre has the best ratings?

7.6 For each of the plots above, note down your observation

### 7.1. Make scatter plot/joinplot for Rating vs. Price

```
sns.jointplot(data=df,y='Rating',x='Price')
```

```
<seaborn.axisgrid.JointGrid at 0x19abbee7400>
```



7.1.1 What pattern do you observe? Does rating increase with price? interpretation: Most of Apps with high price get > 3 Rating but this is because majority of apps are with low price. In addition most apps get rating > 3.


**Concusion:** We cannot consider there is a good relationship between Rating and Price. It seems Price has limited impact on Rating.



## 7.2 Make scatter plot/joinplot for Rating vs. Size

```
sns.jointplot(data=df,y='Rating',x='Size')
```

<seaborn.axisgrid.JointGrid at 0x19ab7cc4a60>



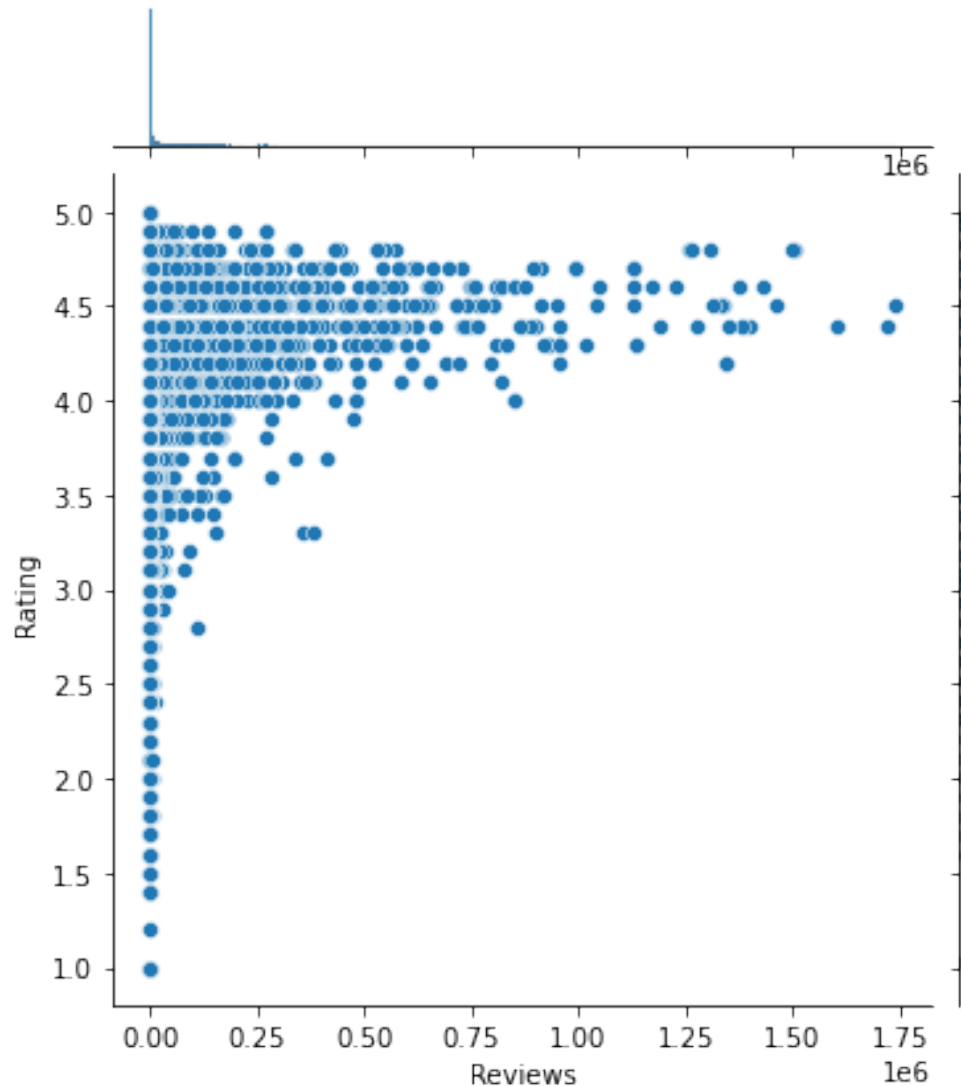
### 7.2.1 Are heavier apps rated better?

**interpretation:**if we look to the area where most apps rated > 3 almost the points are evenly distributed The relationship between Size and rating is very weak

### 7.3. Make scatter plot/joinplot for Rating vs. Reviews

```
sns.jointplot(data=df,y='Rating',x='Reviews')
```

```
<seaborn.axisgrid.JointGrid at 0x19ac176c520>
```



7.3.1 Does more review mean a better rating always?

**interpretation:** Although the relationship seems also not so strong, but we can notice that there is some concentration of apps with higher reviews in high rating area. It seems good apps get more reviews than others


## 7.4 Make boxplot for Rating vs. Content Rating

```
df['Content Rating'].unique()

array(['Everyone', 'Teen', 'Everyone 10+', 'Mature 17+',
      'Adults only 18+', 'Unrated'], dtype=object)

sns.boxplot(data=df, x='Rating', y='Content Rating')

<AxesSubplot:xlabel='Rating', ylabel='Content Rating'>
```




7.4.1 Is there any difference in the ratings? Are some types liked better? Is there any difference in the ratings? Are some types liked better?

**interpretation: Apps of Adults only 18+ has higher rating than others while Mature 17+ gets less likes. Others seem to be closed. Content has good impact on Rating**

## #7.5 Make boxplot for Ratings vs. Category

- `a4_dims = (11.7, 10.27)`  
`fig, ax = plt.subplots(figsize=a4_dims)`  
`sns.boxplot(data=df, x='Rating', y='Category', ax=ax)`

<AxesSubplot:xlabel='Rating', ylabel='Category'>



7.5.1 Which genre has the best ratings?

while observing box plot i found that the Events genere has highest rating.

7.6 For each of the plots above, note down your observation

while observing above plot i founout that Rating vs. Content Rating and Ratings vs. Category. gives better understanding of app rating

## 8. Data preprocessing

For the steps below, create a copy of the dataframe to make all the edits. Name it inp1.

8.1 Reviews and Install have some values that are still relatively very high. Before building a linear regression model, you need to reduce the skew. Apply log transformation (`np.log1p`) to Reviews and Installs.

8.2 Drop columns App, Last Updated, Current Ver, and Android Ver. These variables are not useful for our task.

8.3 Get dummy columns for Category, Genres, and Content Rating. This needs to be done as the models do not understand categorical data, and all data should be numeric. Dummy encoding is one way to convert character fields to numeric. Name of dataframe should be inp2.

### Q 8.1

```
inp1=df.copy()
inp1.Reviews=inp1.Reviews.apply(np.log1p)
```

### Q 8.2

```
inp1.drop(columns=['App','Last Updated','Current Ver','Android Ver'],inplace=True)
```

```
inp1.shape
```

```
(7307, 9)
```

### Q 8.3

- `inp2= pd.get_dummies(inp1)`

```
inp2.shape
```

```
(7307, 158)
```

**9. Train test split and apply 70-30 split. Name the new dataframes df\_train and df\_test.**

**Separate the dataframes into X\_train, y\_train, X\_test, and y\_test.**

- `data = inp2.drop(columns='Rating')`  
`data.shape`

`(7307, 157)`

- `target = pd.DataFrame(inp2.Rating)`  
`target.shape`

`(7307, 1)`

- `x_train, x_test, y_train, y_test = train_test_split(data, target,`  
`test_size=0.3, random_state=3)`  
`print("x_train shape is ", x_train.shape)`  
`print("y_train shape is ", y_train.shape)`  
`print("x_test shape is ", x_test.shape)`  
`print("y_test shape is ", y_test.shape)`

`x_train shape is (5114, 157)`

`y_train shape is (5114, 1)`

`x_test shape is (2193, 157)`

`y_test shape is (2193, 1)`

## 11. Model building

11.1 Use linear regression as the technique

11.2 Report the R2 on the train set

- `model=LinearRegression()`  
`model.fit(x_train, y_train)`
- `LinearRegression()`
- `train_predict=model.predict(x_train)`
- `print("R2 value of the model(by train) is ", r2_score(y_train, train_predict))`

**R2 value of the model(by train) is 0.08010678015666617**

### 1. Make predictions on test set and report R2.

`test_predict=model.predict(x_test)`

- `print("R2 value of the model(by test) is ", r2_score(y_test, test_predict))`

**R2 value of the model(by test) is 0.0522779707476938**

- `pandoc C:\Users\shiwa\Downloads\App Rating Prediction.ipynb -s -o App Rating Prediction.docx`